

# Investigation of some Stochastic Scheduling Problems

Master's Thesis in Computer Science


Philipp Müller

Technische Universität München

November 20, 2013

## Problem statement

- Set of tasks
- Dependencies form intree structure
- Task times exponentially distributed with same parameter
- Nonpreemptive scheduling
- Goal: Minimize total expected make span



# Schedules and snapshots

## Schedules

- Describes the order of tasks
- Our scenario: non-deterministic (task times are random)
- Scheduling *strategy* influences expected make span


## Snapshots

- States of a schedule
- Consist of current intree and the set of scheduled tasks


## Expected run time

- $r$  tasks currently scheduled
- Each of these tasks is the first to finish with probability  $\frac{1}{r}$
- Compute run time recursively (task times memoryless)
  - Expected time for fastest task  $\frac{1}{r}$
  - Weighted expected time for successive snapshots


# Schedule visualization




## Schedule visualization



# Schedule visualization




# Schedule visualization



$$\frac{1}{2} + 0.1 \cdot 6.1640 + 0.4 \cdot 5.9921 + 0.2 \cdot 5.8203 + \dots = 6.43745$$




# Schedule visualization



$$\frac{1}{2} + 0.1 \cdot 6.1640 + 0.4 \cdot 5.9921 + 0.2 \cdot 5.8203 + \dots = 6.43745$$

## Equivalent snapshots

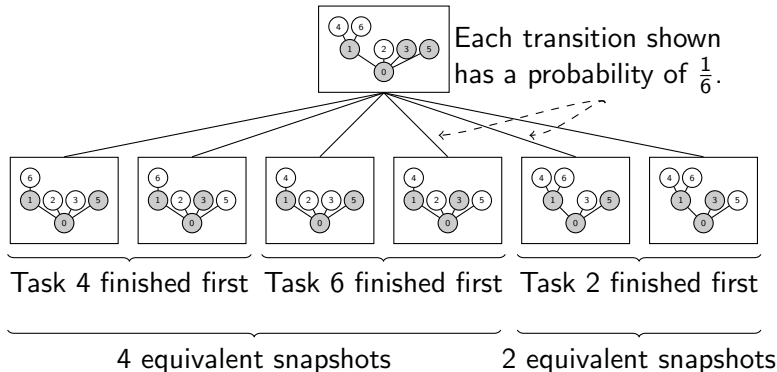
These two snapshots describe the same:




## Equivalent snapshots

- Isomorphic intrees
- Scheduled tasks are mapped onto each other
- Equivalent snapshots excluded/re-used if they result from the same finishing task
- *Canonical* snapshots constructable in  $O(n)$

## Equivalent snapshots — example




# Equivalent snapshots — example



## Two processors — optimal solution

Known results (Chandy, Reynolds 1975)

- Highest-level-first (HLF) optimal for two processors
- Profile (number of tasks per level) completely determines run time
- Reduce computation of expected run time to profiles



## Expected run time — two-leaves intrees

### Theorem (Maaß 2001)

Let  $l, k \in \mathbb{N}$ ,  $a \in \mathbb{N}_0$ . Intrees with profile  $\llbracket (1)^{l-k}, (2)^k, (1)^{a+1} \rrbracket$  have expected run time

$$\begin{aligned} & \sum_{i=1}^k \left(\frac{1}{2}\right)^{l+i-1} \cdot \binom{l+i-2}{i-1} \cdot (k-i+2) \\ & + \sum_{j=1}^l \left(\frac{1}{2}\right)^{k+j-1} \cdot \binom{k+j-2}{j-1} \cdot (l-j+2) \\ & + \sum_{i=1}^k \sum_{j=1}^l \left(\frac{1}{2}\right)^{k-i+l-j+1} \cdot \binom{ki+l-j}{l-j} \\ & + a. \end{aligned}$$

## Expected run time — intrees with many 1-levels

### Theorem

*Intrees with profile  $\llbracket n_1, (1)^{j-2}, n_j, (1)^{r-j} \rrbracket$  (for  $j \geq 2$ ) has expected run time*

$$\mathbb{E} \left[ \llbracket n_1, (1)^{j-2}, n_j, (1)^{r-j} \rrbracket \right] = r + \frac{A_0(n_1 - 2)}{2^{n_1-1}} + \frac{A_{j-1}(n_j - 2)}{2^{n_j+j-2}},$$

where  $A_i$  is inductively defined as follows:

$$A_0(n) = (n + 1) \cdot 2^n \quad A_{i+1}(n) = \sum_{k=0}^n A_i(k)$$



## Profile DAG

- Use profiles instead of whole snapshots
- Scheduled tasks given implicitly (HLF)
- Worst case profile  $\llbracket (1)^{\lfloor \frac{n}{2} \rfloor - 1}, \lceil \frac{n}{2} \rceil, 1 \rrbracket$
- Worst case profile DAG size has  $\lfloor \frac{n}{2} \rfloor \cdot \lceil \frac{n}{2} \rceil + 1$  “profile snapshots”
- Each profile snapshot accounts for less than  $n^2$  “original snapshots”  
⇒ Simple proof: At most  $O(n^4)$  original snapshots.




## HLF may be ambiguous ...

HLF may result in different run times:



... or even strictly suboptimal



## HLF — summary

- HLF is suboptimal . . .
- . . . but asymptotically good (Papadimitriou, Tsitsiklis 1987):  
There is a function  $\beta : \mathbb{N} \mapsto \mathbb{R}_0^+$  with  $\lim_{n \rightarrow \infty} \beta(n) = 0$  such that for each intree  $I$  and an arbitrary HLF strategy  $HLF$  we have

$$T_{HLF}(I) \leq T_{\pi^*}(I) \cdot (1 + \beta(N)),$$

where  $\pi^*$  is the optimal strategy.

- Optimal schedule is often one particular run of HLF **Wie oft – kurze Tabelle in Anhang!**  $\Rightarrow$  HLF is “can-optimal” for these intrees

## (Dynamic) list scheduling

- Can not be optimal for our problem
- Optimal schedule has to consider previous choices
- HLF is particular instance of dynamic list scheduling

## “2-HLF plus 1”



### Motivation

Non-HLF intrees up to 13 tasks have at most one non-HLF task scheduled.

### Strategy

Discard snapshots with more than one non-HLF task scheduled.

### Counterexample





## Subtree with fewer topmost tasks

### Motivation

In many cases, a topmost task being the single requirement for its direct successor has to be scheduled

### Definition (Topmost-maximal subtree for a leaf)


Let  $t$  be a leaf and let  $p = (t, t_1, t_2, t_3, \dots, r)$  be the path from  $t$  to the root  $r$ .

The *topmost-maximal subtree* for  $t$  is the subtree rooted at the *lowest* task  $t^*$  within  $p$  different from  $t$  that does *not* contain more topmost tasks than the subtree rooted at the direct successor of  $t$ .

## Subtree with fewer topmost tasks

Topmost-maximal subtree for task 10:

Example intree






## Subtree with fewer topmost tasks

Topmost-maximal subtree for task 10:


- The subtree rooted at 10 (called  $I_{10}$ ) contains only the topmost task 10 (omitted).

Example intree



## Subtree with fewer topmost tasks

### Example intree




Topmost-maximal subtree for task 10:

- The subtree rooted at 10 (called  $I_{10}$ ) contains only the topmost task 10 (omitted).
- Subtree  $I_8$  contains only topmost task 10 (reference).

## Subtree with fewer topmost tasks

### Example intree




Topmost-maximal subtree for task 10:

- The subtree rooted at 10 (called  $l_{10}$ ) contains only the topmost task 10 (omitted).
- Subtree  $l_8$  contains only topmost task 10 (reference).
- Subtree  $l_3$  still contains only 10 as the topmost task (it introduces only new *leaves*, namely 6 and 7).

## Subtree with fewer topmost tasks

### Example intree




Topmost-maximal subtree for task 10:

- The subtree rooted at 10 (called  $I_{10}$ ) contains only the topmost task 10 (omitted).
- Subtree  $I_8$  contains only topmost task 10 (reference).
- Subtree  $I_3$  still contains only 10 as the topmost task (it introduces only new *leaves*, namely 6 and 7).
- Subtree  $I_1$  contains 10 *and* 11 as topmost tasks.

## Subtree with fewer topmost tasks

### Example intree



Topmost-maximal subtree for task 10:



- The subtree rooted at 10 (called  $I_{10}$ ) contains only the topmost task 10 (omitted).
- Subtree  $I_8$  contains only topmost task 10 (reference).
- Subtree  $I_3$  still contains only 10 as the topmost task (it introduces only new *leaves*, namely 6 and 7).
- Subtree  $I_1$  contains 10 *and* 11 as topmost tasks.
- Subtree  $I_3$  is the topmost maximal subtree for task 10.

## Subtree with fewer topmost tasks

### Strategy

Prefer tasks whose topmost-maximal subtree has fewer topmost tasks.

### Counterexample



## Subtree with fewer leaves

### Motivation

Maybe we were wrong and should have focussed on *ready*, and not on *topmost* tasks.

### Definition (Leaf-maximal subtree for a leaf)

Let  $t$  be a leaf and let  $p = (t, t_1, t_2, t_3, \dots, r)$  be the path from  $t$  to the root  $r$ .

The *leaf-maximal subtree* for  $t$  is the subtree rooted at the *lowest* task  $t^*$  within  $p$  different from  $t$  that does *not* contain more leaves than the subtree rooted at the direct successor of  $t$ .



*Remark:* Preferring subtrees with more does not work.

## Subtree with fewer leaves

### Strategy

Prefer tasks whose leaf-maximal subtree has fewer topmost tasks.

### Counterexample






## Recursive approach

### Strategy

Prefer root's predecessors with maximal processing time.

### Counterexample




## Filling up subtrees

### Motivation

Subtrees seem to be filled up one after another.


### Counterexample




# Maximizing 3-processor time, minimizing 1-processor time

Optimal schedules do *not* necessarily ...


- ... maximize the expected time span  $T_3$  where 3 processors are busy
- ... minimize the expected time span  $T_1$  where only 1 processor can work



$$T_1 = 2.98$$



$$T_3 = 2.36$$



$$T_3 = 2.33, T_1 = 2.99$$

## Optimal strategies — facts


- No deliberate idleness (Chandy, Reynolds 1979)
- More processors do not worsen run time (Maaß 2001)
- Preemptive scheduling is strictly better than non-preemptive scheduling for 3 or more processors (equal for 2 processors)
- “Optimal schedules for subtrees do not help”

## Computing optimal schedules


- Exhaustive search (*LEAF* scheduler)
- “Optimize” result from exhaustive search by recursively discarding bad choices
- Exponential run time

## Size of the snapshot DAG

- Number of subtrees grows exponentially  
⇒ Number of overall snapshots as well
- Number of snapshots in optimal DAG remarkably smaller



Maximum



Average

## Compressing the snapshot DAG

- Optimal snapshot DAG allows (recursive) merging of some snapshots
- LEAF snapshot DAG prevents us from merging because the run times are not the same for some snapshots **Verbessern!**

## Optimal strategies — conjectures

- *In the beginning*, as many topmost task as possible shall be scheduled
- If the scheduler has a choice, it shall pick a topmost task, if available
- If only non-top tasks are scheduled, we can exchange any one with a topmost task and can obtain a better run time



## Degenerate intrees

- On each level, at most one task has predecessors
- Uniquely determined by their profile

### Theorem

*Degenerate intrees are optimally scheduled by HLF.*

### Proof.

- Assert that for degenerate intree  $I$ , two tasks  $z_1, z_2$  with  $level(z_1) \geq level(z_2)$ :  $T_{x,y,z_1}^*(I \cup \{z_1\}) \geq T_{x,y,z_2}^*(I \cup \{z_2\})$
- Induction over the number of tasks, comparing  $T_{HLF}(I)$  to  $T_{x',y',z'}(I)$




## Parallel chains

- Each task, except the root, has at most one predecessor
- Parallel chains with up to 27 tasks are optimally scheduled by HLF

## Degenerate intrees and parallel chains — outcome

- HLF is deterministic for these classes ...
- ... and needs remarkably less snapshots than LEAF
- Example: Number of snapshots for degenerate binary trees




- Analogous results for parallel chains

# Improving LEAF with conjectures

## SCLEAF

- Simple LEAF scheduler, but ...
- ... using HLF when encountering degenerate intrees or parallel chains ...
- ... restricts to snapshots where as many topmost task as possible are scheduled



## Practical results

- Excluding equivalent snapshots speeds up the program by at least factor 3 (for intrees with 11 or more tasks)
- Computing optimal schedules for all intrees with up to 15 tasks in 11 minutes and less than 2Gb of main memory
- More tasks require too much memory
- Computing optimal schedules for non-trivial intrees with 19 tasks takes one day
- Using Boost Rational to represent expectancies as fractions requires slightly more time and memory
- Using GMP to represent expectancies as fractions requires roughly doubles time and increases memory consumption by roughly 40%